Membership Inference Auditing of Small SMILES Language Models on MOSES: LiRA, Randomized-SMILES Training, and Preliminary DP-SGD Results

Ankit Dhattarwal

Abstract—Autoregressive models trained on molecular SMILES strings are used to explore generative chemistry; some drug-discovery models may be trained on confidential compound collections. In this study, we empirically evaluate Membership Inference Attacks (MIA) against small causal SMILES language models (2 layers, $d = 128$, with 419,584 parameters, vocab size 25) on a subset of the MOSES benchmark across training epochs (5, 20, 60), dataset sizes ($N \in \{1\text{k}, 2.5\text{k}, 5\text{k}, 10\text{k}\}$), and defenses. Under canonical SMILES queries and sequence negative log-likelihood sum (*nll_sum*), models exhibit substantial membership leakage by epoch 60: online LiRA reaches a ROC-AUC of 0.898 ± 0.003 and a True Positive Rate (TPR) of $27.82 \pm 1.47\%$ at a 1% False Positive Rate (FPR), compared to 0.689 ± 0.005 AUC and $2.96 \pm 0.24\%$ TPR@1% FPR for standard LOSS. In smaller training cohorts ($N = 1,000$), attack sensitivity increases: LOSS AUC reaches 0.998 ± 0.001 (TPR@1% $= 93.99 \pm 4.35\%$) and online LiRA AUC reaches 0.975 ± 0.004 (TPR@1% $= 45.49 \pm 33.14\%$).
Mean-aggregated scoring over up to 10 string forms (canonical included) removes most of the signal (online LiRA AUC drops to 0.528 ± 0.006); this reflects a change in the query/scoring method for a model trained on canonical SMILES and is not a defense. Finally, in our preliminary DP-SGD arm (target $\epsilon \in \{8.0, 3.0\}$, $\delta = 10^{-5}$), generative utility was not evaluated, so the low attack AUC ($0.500 - 0.528$) cannot be interpreted as an effective defense. Training with randomized SMILES reduces canonical-query offline LiRA AUC to 0.528 ± 0.004 (variants scoring: 0.552 ± 0.005); validation NLL and generative validity of these models were not evaluated. Of 15,000 generated samples, 12 canonicalized outputs matched molecules in the training set and 7 matched molecules in the holdout set. Conditioning on these 19 matches, assuming equal match probability for the two equally sized reference sets, and treating the 19 match outcomes as independent under H0 (p = 0.5), the one-sided exact binomial p-value is 0.1796 and the two-sided exact binomial p-value is 0.3593. Neither result is significant at the 0.05 level.

Index Terms—Membership Inference Attacks, Molecular Generative Models, SMILES, Differential Privacy, Likelihood-Ratio Attack, Privacy Auditing.

1 Introduction

Autoregressive generative sequence models parameterize discrete token distributions to explore chemical space. By modeling the Simplified Molecular Input Line Entry System (SMILES) [1], causal transformers can generate candidate structures and propose chemical variations for lead optimization.

Some drug-discovery models are trained or fine-tuned on confidential compound collections. If a model is shared, an adversary with likelihood access may use membership inference to assess whether a target compound occurred in its training data, creating a potential confidentiality risk [2], [3].

While membership inference has been widely investigated in computer vision and natural language processing, auditing molecular generative sequence models presents distinctive factors:

1) Molecules admit multiple non-unique equivalent SMILES representations (non-canonical string enumeration).
2) Chemical libraries often contain recurring core frameworks (Bemis-Murcko scaffolds [4]) that induce structural correlation across compounds.
3) Language models may memorize specific surface syntax without learning generalized semantic invariants.

In this work, we present an empirical audit of membership leakage in small SMILES autoregressive models. Using a shadow-model ensemble framework based on Carlini et al.'s Likelihood-Ratio Attack (LiRA) [5], we evaluate exposure across training duration, dataset scaling, transfer learning, and defensive interventions.

2 Threat Model and Formulation

2.1 Sequence Modeling and Scoring Metric

Let a molecular SMILES string be represented as a token sequence $x = (x_1, \ldots, x_T) \in \mathcal{V}^T$, where $\mathcal{V}$ is a discrete vocabulary of chemical tokens. An autoregressive transformer parameterizes the likelihood:

$$P_\theta(x) = \prod_{t=1}^{T} P_\theta(x_t \mid x_{<t}) \tag{1}$$

A. Dhattarwal is an independent researcher in Machine Learning and Computational Chemistry. Project code and experiment artifacts are available at: https://github.com/Ankit-Dhattarwal/chem-generative-privacy-audit

The sequence negative log-likelihood (NLL) sum is defined as:

$$\mathrm{NLL}_{\mathrm{sum}}(\theta; x) = -\sum_{t=1}^{T} \log P_\theta(x_t \mid x_{<t}) \tag{2}$$

Throughout this paper, we consistently report and evaluate metrics using sequence NLL sum (*nll_sum*). In our experiments, *nll_sum* and *nll_mean* yield near-identical LiRA results across epochs 5, 20, and 60 (LiRA z-scores are mathematically invariant to per-molecule length scaling), while uncalibrated LOSS differs only marginally ($\Delta\mathrm{AUC} \leq 0.004$); we report *nll_sum* consistently throughout. Where indicated, we also report randomized-SMILES variant scoring (*nll_variants_mean*), which averages NLL across up to 10 distinct SMILES strings per molecule (the canonical form plus randomized forms), under mean aggregation (individual non-canonical forms were not scored separately).

2.2 Attack Formulations: LOSS vs. LiRA

Given a query molecule x, the adversary computes a test statistic $S(x)$ to predict whether $x \in \mathcal{D}_{\mathrm{train}}$:

1) Loss Thresholding (LOSS):

$$S_{\mathrm{LOSS}}(x) = -\mathrm{NLL}_{\mathrm{sum}}(\mathcal{M}_{\mathrm{target}}; x) \tag{3}$$

LOSS assumes that training examples achieve lower sequence cross-entropy loss than non-training examples. However, intrinsic molecular complexity varies substantially: complex macrocycles or long branched chains exhibit higher baseline loss than simple aromatics regardless of membership.

2) Calibrated Likelihood-Ratio Attack (LiRA): To adjust for molecular complexity, LiRA [5] trains an ensemble of K shadow models $\{\mathcal{M}_k\}_{k=1}^{K}$ on random subsets of the data pool. For each query x, the models are partitioned into in-models $\mathcal{S}_{\mathrm{in}}(x) = \{k : x \in \mathcal{D}_k\}$ and out-models $\mathcal{S}_{\mathrm{out}}(x) = \{k : x \notin \mathcal{D}_k\}$.

- Offline LiRA (Out-only): Estimates the out-distribution parameters $\mu_{\mathrm{out}}(x)$ and $\sigma_{\mathrm{out}}(x)$ over out-models:

$$\Lambda_{\mathrm{offline}}(x) = -\frac{\mathrm{NLL}_{\mathrm{sum}}(\mathcal{M}_{\mathrm{target}}; x) - \mu_{\mathrm{out}}(x)}{\sigma_{\mathrm{out}}(x) + 10^{-6}} \tag{4}$$

- Online LiRA (In vs. Out): Estimates both in-distribution $(\mu_{\mathrm{in}}, \sigma_{\mathrm{in}})$ and out-distribution $(\mu_{\mathrm{out}}, \sigma_{\mathrm{out}})$ parameters, computing the log-likelihood ratio under parametric Gaussian assumptions:

$$\begin{aligned}\Lambda_{\mathrm{online}}(x) = &-\frac{1}{2}\left(\frac{\mathrm{obs} - \mu_{\mathrm{in}}}{\sigma_{\mathrm{in}}}\right)^2 \\ &+\frac{1}{2}\left(\frac{\mathrm{obs} - \mu_{\mathrm{out}}}{\sigma_{\mathrm{out}}}\right)^2 \\ &+\log(\sigma_{\mathrm{out}}/\sigma_{\mathrm{in}})\end{aligned} \tag{5}$$

2.3 Threat Model Assumptions

We evaluate our models under an explicit, optimistic threat model:

1) Distribution Match: Shadow models are sampled from the exact same marginal distribution as the target model (MOSES benchmark pool).
2) Likelihood Access: The adversary has query access enabling exact calculation of token-level sequence log-likelihoods (*nll_sum*).
3) Canonical-String Queries: The adversary queries the target and shadow models using standard canonical SMILES representations.

These represent strong assumptions favoring the attacker. In restricted black-box deployments (e.g., API rate limits, generation-only access, or query filtering), adversary effectiveness may be substantially lower.

3 Experimental Setup and Architecture

3.1 Model Architecture and Training Setup

All evaluations utilize a small causal GPT architecture: 2 transformer layers, hidden embedding dimension $d_{\mathrm{model}} = 128$, 4 attention heads, dropout 0.1, and a context window of 128 tokens. Instantiating the model over the dataset tokenizer (len(tokenizer.itos)) yields an exact vocabulary size of 25 tokens and 419,584 parameters. Models are trained using AdamW (lr $= 10^{-3}$, weight_decay $= 0.01$), batch size 32, with gradient clipping norm 1.0 and a fixed sequence length of 128 tokens. Differentially private models (DPCausalGPT) evaluate the identical architecture (419,584 parameters) and utilize the same optimizer configuration (AdamW, lr $= 10^{-3}$, weight_decay $= 0.01$, batch size 32, dropout 0.1, max_grad_norm $= 1.0$) wrapped via Opacus.

3.2 Dataset and Shadow Model Partitioning

Data are sourced from the MOSES benchmark [6]. The evaluation pool size is POOL_N $= 20,000$ molecules. All experiments use the public MOSES benchmark; the terms member and private refer to an experimental split, not confidential real-world records. For each shadow model k, exactly $N = 10,000$ molecules are sampled uniformly at random without replacement using seed k (_membership(pool, k) in colab_pipeline.py). Evaluation is conducted across all 20,000 candidate compounds (10,000 members, 10,000 non-members). In the baseline and randomized-augmentation (E5a) experiments, 32 shadow models are trained ($k \in \{0, \ldots, 31\}$). In data scaling (E3), DP-SGD, and pretrain-finetune experiments, 8 shadow models are trained per condition.

3.3 Opacus DP-SGD Implementation

Differentially private training is implemented using Opacus (version 1.6.0) [7]. The noise scale σ is computed dynamically using opacus.accountants.utils.get_noise_multiplier(target_epsilon=eps,target_delta=1e-5,sample_rate=32/N,epochs=60). Gradient clipping is set to max_grad_norm $= 1.0$, targeting $\epsilon \in \{8.0, 3.0\}$ with $\delta = 10^{-5}$. The final saved accountant values were $\epsilon = 7.158$ for target 8 and $\epsilon = 2.750$ for target 3 across the eight models per setting. The causal attention layer avoids fused operations incompatible with Opacus per-sample gradient hooks. If training resumes from a checkpoint, privacy accounting restarts from zero unless serialized accountant state is restored.

4 Empirical Results and Findings

4.1 Baseline Leakage vs. Training Duration

Table 1 and Fig. 1 present membership inference performance across 5, 20, and 60 training epochs on canonical SMILES representations evaluated using *nll_sum*.

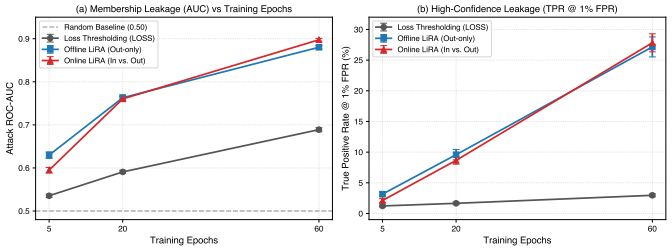


Fig. 1. Baseline Attack Vulnerability vs. Training Epochs. (a) ROC-AUC progression. (b) True Positive Rate at 1% False Positive Rate.

At epoch 5, membership leakage is modest (strongest attack: offline LiRA AUC 0.630 ± 0.008). By epoch 60, extensive exposure occurs: online LiRA reaches ROC-AUC 0.898 ± 0.003, with TPR@1% FPR reaching $27.82 \pm 1.47\%$. At a stricter threshold of 0.1% FPR, online LiRA still identifies $8.15 \pm 1.58\%$ of member molecules. Standard loss thresholding substantially underestimates vulnerability (AUC 0.689 vs. 0.898 for online LiRA at epoch 60).

Statistical Uncertainty: The reported $\pm$ values reflect the standard deviation across 32 leave-one-out shadow target models. To assess uncertainty over the 20,000 evaluation molecules, we computed 95% bootstrap confidence intervals for epoch-60 online LiRA across three target models (not independent: they share shadow models and overlapping training pools):

- Target 0: AUC 0.9004, 95% CI $[0.8962, 0.9045]$; TPR@1% 0.2684, 95% CI $[0.2492, 0.3000]$.
- Target 1: AUC 0.8916, 95% CI $[0.8870, 0.8961]$; TPR@1% 0.2510, 95% CI $[0.2178, 0.2756]$.
- Target 2: AUC 0.8972, 95% CI $[0.8930, 0.9012]$; TPR@1% 0.2838, 95% CI $[0.2613, 0.3136]$.

4.2 String-Form Sensitivity: Randomized-SMILES Scoring

In our setting, mean-aggregated scoring over up to 10 string forms (canonical included) removes most of the signal. In colab_pipeline.py, each molecule was scored at epoch 60 using up to 10 distinct SMILES strings per molecule (the canonical form plus randomized forms generated via RDKit Chem.MolToSmiles(mol, doRandom=True, canonical=False)), with the average NLL recorded as *nll_variants_mean*.

As reported in Table 1 (bottom section), scoring canonical models with randomized variants drops LOSS AUC from 0.689 to 0.511 ± 0.004, offline LiRA AUC from 0.880 to 0.560 ± 0.007, and online LiRA AUC from 0.898 to 0.528 ± 0.006. We did not score individual non-canonical strings, so we cannot say whether single randomized forms leak.

4.3 Private Dataset Scaling

We examine vulnerability across dataset sizes $N \in \{1\text{k}, 2.5\text{k}, 5\text{k}, 10\text{k}\}$ trained for 60 epochs, with 8 shadow models per size ($2N$ candidate universe). Results are presented in Table 2 and Fig. 2, separating LOSS and LiRA attacks.

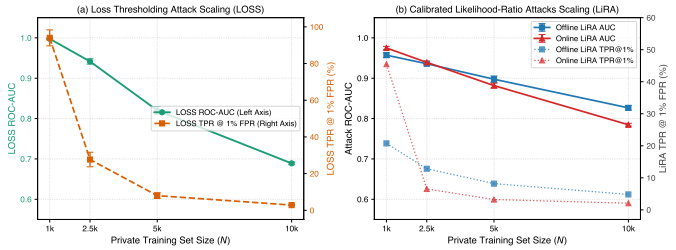


Fig. 2. Dataset Size Scaling (E3). (a) Loss thresholding attack (LOSS). (b) Calibrated Likelihood-Ratio Attacks (LiRA).

Smaller training cohorts experience pronounced overfitting at 60 epochs. At $N = 1,000$, standard LOSS reaches AUC 0.998 ± 0.001 with TPR@1% $= 93.99 \pm 4.35\%$, while offline LiRA reaches AUC 0.957 ± 0.006 (TPR@1% $= 20.75 \pm 6.51\%$) and online LiRA reaches AUC 0.975 ± 0.004 (TPR@1% $= 45.49 \pm 33.14\%$). With only 8 shadow models, LiRA density estimation at small N suffers from high variance (e.g., online LiRA TPR@1% standard deviation is $\pm 33.14\%$). We report these empirical numbers directly and do not claim a formal parametric scaling law.

4.4 Transfer Learning: Pretraining to Fine-Tuning

We evaluated a transfer learning setting where a model was pretrained on 10,000 public molecules for 60 epochs, then fine-tuned on 10,000 private molecules for 20 epochs across 8 shadow targets (Table 3).

Offline LiRA achieves an AUC of 0.801 ± 0.003 and TPR@1% of $4.36 \pm 0.51\%$. While lower than baseline epoch 60 (0.880 ± 0.004 / $27.15 \pm 1.63\%$), offline LiRA AUC is higher than the baseline trained from scratch for 20 epochs (0.763 ± 0.007 / $9.58 \pm 0.84\%$, though TPR@1% is lower), indicating that fine-tuning on private data remains susceptible to membership inference. However, there is no matched from-scratch control trained on the identical 10,000 private compounds for 20 epochs without pretraining, so the isolated effect of pretraining cannot be disentangled.

4.5 Defense Comparison at Epoch 60

Table 4 and Fig. 3 compare two defensive approaches at epoch 60: (1) Randomized-SMILES training (E5a), where models are trained with random-SMILES data augmentation applied per batch; and (2) DP-SGD targeting $\epsilon \in \{8.0, 3.0\}$ with $\delta = 10^{-5}$.

Randomized-SMILES training reduces canonical attack performance: offline LiRA drops from 0.880 to 0.528 ± 0.004, and TPR@1% drops from 27.15% to 1.38%. When attacked via randomized variants scoring, offline LiRA AUC remains low at 0.552 ± 0.005.

In preliminary runs, DP-SGD reduces attack AUC near random guess ($0.500 - 0.528$). However, generative utility was not evaluated: an attack AUC near 0.5 can arise from an under-trained model rather than an effective defense,

TABLE 1
Baseline membership inference vulnerability across training epochs on canonical SMILES. Evaluated over 32 shadow models on 20,000 candidate targets (10,000 members, 10,000 non-members). Values report mean ± standard deviation across 32 leave-one-out target evaluations (not a confidence interval). All metrics use negative log-likelihood sum (nll_sum), with randomized-SMILES variants scoring (up to 10 distinct SMILES strings per molecule, the canonical form plus randomized forms, mean NLL) reported at epoch 60.

Epochs	Scoring Representation	Attack Method	ROC-AUC	TPR @ 1% FPR (%)	TPR @ 0.1% FPR (%)
5	Canonical (nll_sum)	Loss Thresholding (LOSS)	0.535 ± 0.004	1.23 ± 0.18	0.14 ± 0.07
5	Canonical (nll_sum)	Offline LiRA (Out-only)	0.630 ± 0.008	3.15 ± 0.36	0.42 ± 0.15
5	Canonical (nll_sum)	Online LiRA (In vs. Out)	0.595 ± 0.006	2.12 ± 0.35	0.19 ± 0.09
5	Canonical (nll_sum)	Strongest: offline_LiRA	0.630 ± 0.008	3.15 ± 0.36	0.42 ± 0.15
20	Canonical (nll_sum)	Loss Thresholding (LOSS)	0.591 ± 0.004	1.66 ± 0.21	0.21 ± 0.08
20	Canonical (nll_sum)	Offline LiRA (Out-only)	0.763 ± 0.007	9.58 ± 0.84	1.86 ± 0.47
20	Canonical (nll_sum)	Online LiRA (In vs. Out)	0.760 ± 0.006	8.66 ± 0.62	1.34 ± 0.33
20	Canonical (nll_sum)	Strongest: offline_LiRA	0.763 ± 0.007	9.58 ± 0.84	1.86 ± 0.47
60	Canonical (nll_sum)	Loss Thresholding (LOSS)	0.689 ± 0.005	2.96 ± 0.24	0.41 ± 0.15
60	Canonical (nll_sum)	Offline LiRA (Out-only)	0.880 ± 0.004	27.15 ± 1.63	8.22 ± 1.32
60	Canonical (nll_sum)	Online LiRA (In vs. Out)	0.898 ± 0.003	27.82 ± 1.47	8.15 ± 1.58
60	Canonical (nll_sum)	Strongest: online_LiRA	0.898 ± 0.003	27.82 ± 1.47	8.15 ± 1.58
60 (variants)	Randomized-SMILES (up to 10 strings, mean NLL)	Loss Thresholding (LOSS)	0.511 ± 0.004	1.09 ± 0.12	0.12 ± 0.05
60 (variants)	Randomized-SMILES (up to 10 strings, mean NLL)	Offline LiRA (Out-only)	0.560 ± 0.007	2.01 ± 0.29	0.30 ± 0.10
60 (variants)	Randomized-SMILES (up to 10 strings, mean NLL)	Online LiRA (In vs. Out)	0.528 ± 0.006	1.45 ± 0.47	0.18 ± 0.12
60 (variants)	Randomized-SMILES (up to 10 strings, mean NLL)	Strongest: offline_LiRA	0.560 ± 0.007	2.01 ± 0.29	0.30 ± 0.10

TABLE 2
Membership leakage scaling as private dataset size varies ($N \in \{1k, 2.5k, 5k, 10k\}$). Each size uses 8 shadow models ($2N$ candidate universe). LiRA estimates exhibit high variance at small N due to limited shadow samples per molecule. Evaluated using nll_sum. We present empirical measurements without assuming a parametric power law.

	Loss Thresholding (LOSS)		Offline LiRA		Online LiRA	
Size (N)	ROC-AUC	TPR@1% (%)	ROC-AUC	TPR@1% (%)	ROC-AUC	TPR@1% (%)
1,000	0.998 ± 0.001	93.99 ± 4.35	0.957 ± 0.006	20.75 ± 6.51	0.975 ± 0.004	45.49 ± 33.14
2,500	0.941 ± 0.006	27.68 ± 3.97	0.936 ± 0.003	12.85 ± 2.97	0.939 ± 0.003	6.53 ± 1.60
5,000	0.820 ± 0.010	7.99 ± 1.55	0.898 ± 0.008	8.19 ± 1.05	0.882 ± 0.006	3.23 ± 0.60
10,000	0.689 ± 0.003	2.84 ± 0.09	0.827 ± 0.006	4.81 ± 0.63	0.785 ± 0.003	2.07 ± 0.28

TABLE 3
Membership leakage in transfer learning (public pretraining on 10,000 molecules, fine-tuning on 10,000 private molecules). Metrics report mean ± std across 8 target models (not independent: they share shadow models and overlapping training pools). Note: there is no from-scratch control trained at the exact same size/epochs in this experiment.

Attack Method	ROC-AUC	TPR @ 1% FPR (%)	TPR @ 0.1% FPR (%)
Loss Thresholding (LOSS)	0.640 ± 0.003	2.31 ± 0.26	0.34 ± 0.08
Offline LiRA	0.801 ± 0.003	4.36 ± 0.51	0.43 ± 0.26
Online LiRA	0.751 ± 0.004	2.00 ± 0.22	0.22 ± 0.05

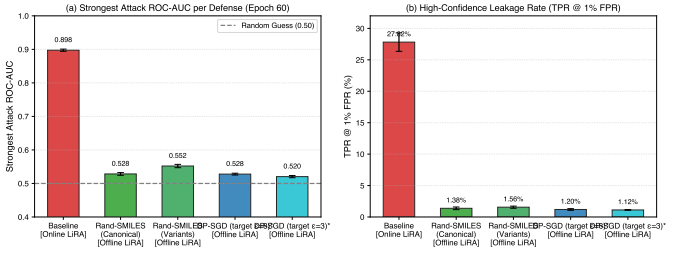


Fig. 3. Defense Comparison at Epoch 60 showing the strongest attack per defense. *Note: preliminary, utility not measured.

and privacy-utility tradeoffs cannot be established without verifying chemical validity. This represents a methodological caution rather than a measured finding of this paper.

4.6 Generative Memorization Audit

To audit whether models generate molecules matching the canonicalized structures in the training set, we sampled 5,000 molecules per model for three baseline models (Models 0, 15, 31; total 15,000 samples) on Apple Silicon (MPS). Generated molecules were evaluated for chemical validity, uniqueness, and canonicalized structure matches against 10,000 training members and 10,000 holdout molecules (Table 5 and Fig. 4).

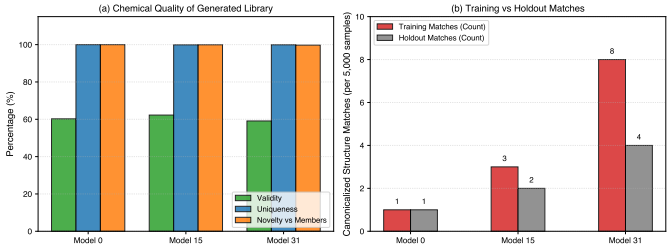


Fig. 4. Generative Memorization and Chemical Quality Audit (E6) across 15,000 samples.

Of 15,000 generated samples, 12 canonicalized outputs matched molecules in the training set and 7 matched molecules in the holdout set. Conditioning on these 19 matches, assuming equal match probability for the two equally sized reference sets, and treating the 19 match outcomes as independent under H0 (p = 0.5), the one-sided exact binomial p-value is 0.1796 and the two-sided exact binomial

TABLE 4
Comparison of privacy defenses at epoch 60 under standard and calibrated attacks. All values use nll_sum (except randomized variants scoring which evaluates mean NLL across up to 10 distinct SMILES strings per molecule). The strongest attack per defense is bolded. *Note: preliminary, utility not measured.

Panel (a): Calibrated Likelihood-Ratio Attacks (LiRA)

		Offline LiRA		Online LiRA		Strongest
Defense Configuration	Shadows	AUC	TPR@1% (%)	AUC	TPR@1% (%)	Attack
Baseline (Unmitigated)	32	0.880 ± 0.004	27.15 ± 1.63	0.898 ± 0.003	27.82 ± 1.47	Online LiRA
Rand-SMILES (Canonical)	32	0.528 ± 0.004	1.38 ± 0.21	0.508 ± 0.004	1.09 ± 0.17	Offline LiRA
Rand-SMILES (Variants, ≤ 10 strings)	32	0.552 ± 0.005	1.56 ± 0.17	0.519 ± 0.005	1.14 ± 0.18	Offline LiRA
DP-SGD (target $\epsilon = 8.0$, $\delta = 10^{-5}$)*	8; prelim.	0.528 ± 0.003	1.20 ± 0.14	0.501 ± 0.005	1.09 ± 0.16	Offline LiRA
DP-SGD (target $\epsilon = 3.0$, $\delta = 10^{-5}$)*	8; prelim.	0.520 ± 0.003	1.12 ± 0.06	0.500 ± 0.002	1.01 ± 0.10	Offline LiRA

Panel (b): Uncalibrated Loss Thresholding Baseline

Defense Configuration	Shadows	LOSS AUC	LOSS TPR@1% (%)
Baseline (Unmitigated)	32	0.689 ± 0.005	2.96 ± 0.24
Rand-SMILES (Canonical)	32	0.506 ± 0.004	1.08 ± 0.16
Rand-SMILES (Variants, ≤ 10 strings)	32	0.506 ± 0.004	1.10 ± 0.16
DP-SGD (target $\epsilon = 8.0$, $\delta = 10^{-5}$)*	8; prelim.	0.505 ± 0.004	1.13 ± 0.15
DP-SGD (target $\epsilon = 3.0$, $\delta = 10^{-5}$)*	8; prelim.	0.503 ± 0.004	1.10 ± 0.06

TABLE 5
Generative memorization audit (E6), with 5,000 samples per model. Counts are canonicalized structure matches. Quality metrics are unweighted means of per-model rates; match counts are pooled. Conditional on 19 matches (12 training, 7 holdout), exact binomial tests under $H_0 : p = 0.5$ give one-sided $p = 0.1796$ and two-sided $p = 0.3593$; neither is significant at 0.05.

Model ID	Validity (%)	Uniqueness (%)	Novelty (%)	Training Matches	Holdout Matches	Binomial p-value	Statistically Significant?
Model 0	60.24%	99.97%	99.97%	1	1	0.7500	No ($p > 0.05$)
Model 15	62.26%	99.87%	99.90%	3	2	0.5000	No ($p > 0.05$)
Model 31	59.10%	99.90%	99.73%	8	4	0.1938	No ($p > 0.05$)
Mean rates; pooled hits	60.53%	99.91%	99.87%	12	7	0.1796	No ($p > 0.05$)

p-value is 0.3593. Neither result is significant at the 0.05 level. Individual model one-sided binomial p-values are 0.750, 0.500, and 0.194.

5 Related Work

Our work interfaces with privacy auditing, generative chemistry, and membership inference:

- Membership Inference Attacks: Carlini et al. [5] established Likelihood-Ratio Attacks (LiRA) to overcome sequence hardness confounding. Shokri et al. [3] originally formulated shadow-model attacks. Recent theoretical analysis by Zhang et al. [8] emphasized that empirical membership inference attacks face significant evidentiary caveats and cannot formally prove that a specific sample was present in a training set without known reference baselines.
- Privacy in Cheminformatics & Molecular Models: Krüger et al. [2] investigated membership leakage in molecular property predictors (molprivacy), demonstrating that publishing neural network weights poses data confidentiality risks. Aubinais et al. [9] proposed an attack-independent indicator for quantized models and evaluated it on molecular modeling tasks. Ross et al. [10] developed GP-MoLFormer, pretrained on a corpus of 1.1 billion SMILES strings, including 650 million unique strings.
- Privacy Auditing in Biological Sequence Models: Nemecek et al. [11] evaluated genomic language models using perplexity-based detection, canary extraction, and membership inference, illustrating domain-specific privacy risks.
- SMILES Randomization & Data Augmentation: Arús-Pous et al. [12] and Bjerrum [13] demonstrated that randomized SMILES enumeration expands the effective representation space and regularizes generative models against overfitting.
- Differential Privacy in Deep Learning: Abadi et al. [14] and the Opacus library [7] formalized differential privacy via DP-SGD and Rényi Differential Privacy (RDP) accounting.

Our contribution provides a focused empirical audit of small causal SMILES language models across training epochs, dataset sizes, randomized representation scoring, and DP-SGD training. We do not claim to present the first or definitive study, but rather a transparent, reproducible empirical audit of these dynamics in small chemical generative LMs.

6 Limitations

This preliminary study has several explicit limitations:

1) Model Capacity and Architecture: Models evaluate a 2-layer, $d = 128$, ~ 0.4M parameter GPT with modest chemical validity ($\sim 60\%$). These findings may not generalize directly to larger generative models.
2) Dataset Scope: Experiments are restricted to a single filtered subset of the public MOSES benchmark. The

study does not evaluate proprietary or confidential collections; diverse private libraries with differing scaffold distributions may exhibit different leakage profiles.

3) Shadow Ensemble Scale Asymmetry: While baseline and E5a arms used 32 shadow models, compute limits restricted E3, DP-SGD, and pretrain-FT to 8 shadow models, leading to higher estimation variance.
4) Reported Uncertainty: Metrics report standard deviation across leave-one-out shadow target models, rather than per-molecule bootstrap confidence intervals. Our computed 95% bootstrap CIs across molecules for epoch 60 online LiRA on three targets span roughly ± 0.005 in AUC and ± 0.03 in TPR@1% (raw per-molecule scoring parquets are omitted from the release repository due to storage constraints and are available on request).
5) Generative Sample Size and Hardware: Experiment E6 sampled 5,000 molecules per model (15,000 total) and ran on Apple Silicon (MPS), whereas shadow training was conducted on CUDA Tesla T4. Statistical tests do not show a significant difference at this sample size ($p > 0.05$).
6) DP-SGD Generative Utility: The DP-SGD arm represents a preliminary exploration where utility was not evaluated for DP checkpoints; an attack AUC near 0.5 may reflect an under-trained model rather than an effective defense. The recorded accountant totals are $\epsilon = 7.158$ for target 8 and $\epsilon = 2.750$ for target 3 at $\delta = 10^{-5}$. This checks the saved accountant totals, not the end-to-end correctness of the training implementation or its assumptions.
7) Lack of From-Scratch Control: The pretrain $\rightarrow$ finetune experiment lacks an exact from-scratch baseline trained for 20 epochs on 10,000 molecules, preventing full disentanglement of pretraining effects from training duration.
8) Randomized-SMILES Utility Evaluation: Validation NLL and generative chemical validity of randomized-SMILES-trained models were not assessed.

7 Conclusion

This preliminary study audited membership leakage in small SMILES language models. We observed that under canonical queries and likelihood access, models overfit over 60 epochs, enabling calibrated LiRA attacks to reach an AUC of 0.898. In our setting, scoring with mean NLL over up to 10 string forms (canonical included) removes most of the signal; training with randomized SMILES also reduces attack AUC to $0.51-0.56$. While DP-SGD is designed to bound the influence of individual training examples, preliminary implementations exhibited possible loss of generative utility (not evaluated); an attack AUC near 0.5 may reflect an under-trained model and is not evidence of a useful defense. Future work must investigate scaling laws on larger architectures, higher-power generative sampling, and privacy-preserving training that maintains chemical validity.

Acknowledgments

AI coding and writing assistants were used during code development, table generation and manuscript editing. The author designed the experiments, reviewed all code, results and text, and takes full responsibility for the content.

References

[1] D. Weininger, "Smiles, a chemical language and information system. 1. introduction to methodology and encoding rules," Journal of Chemical Information and Computer Sciences, vol. 28, no. 1, pp. 31–36, 1988.
[2] F. P. Krüger, J. Östman, L. Mervin, I. V. Tetko, and O. Engkvist, "Publishing neural networks in drug discovery might compromise training data privacy," Journal of Cheminformatics, vol. 17, no. 1, p. 38, 2025.
[3] R. Shokri, M. Stronati, C. Song, and V. Shmatikov, "Membership inference attacks against machine learning models," IEEE Symposium on Security and Privacy (S&P), pp. 3–18, 2017.
[4] G. W. Bemis and M. A. Murcko, "The properties of known drugs. 1. molecular frameworks," Journal of Medicinal Chemistry, vol. 39, no. 15, pp. 2887–2893, 1996.
[5] N. Carlini, S. Chien, M. Nasr, S. Song, A. Terzis, and F. Tramer, "Membership inference attacks from first principles," IEEE Symposium on Security and Privacy (S&P), pp. 1897–1914, 2022.
[6] D. Polykovskiy, A. Zhebrak, B. Sanchez-Lengeling, S. Golovanov, O. Tatanov, S. Belyaev, R. Kurbanov, A. Artamonov, V. Aladinskiy, M. Veselov, A. Kadurin, S. Johansson, H. Chen, S. Nikolenko, A. Aspuru-Guzik, and A. Zhavoronkov, "Molecular sets (moses): A benchmarking platform for molecular generation models," Frontiers in Pharmacology, vol. 11, p. 565644, 2020.
[7] A. Yousefpour, I. Shintre, A. Serrano, P. Overman, S. Zhao, H. Yang, H. Zou, N. Zhang, C. Biron, S. Gopi et al., "Opacus: User-friendly differential privacy library in pytorch," arXiv preprint arXiv:2109.12298, 2021.
[8] J. Zhang, D. Das, G. Kamath, and F. Tramèr, "Membership inference attacks cannot prove that a model was trained on your data," arXiv preprint arXiv:2409.19798, 2024.
[9] E. Aubinais, P. Formont, P. Piantanida, and E. Gassiat, "An indicator of membership inference security in post-training quantized models," in Proceedings of the 29th International Conference on Artificial Intelligence and Statistics, ser. Proceedings of Machine Learning Research, vol. 300. PMLR, 2026, pp. 649–657. [Online]. Available: https://proceedings.mlr.press/v300/aubinais26a.html
[10] J. Ross, B. Belgodere, S. C. Hoffman, V. Chenthamarakshan, J. Navratil, Y. Mroueh, and P. Das, "Gp-molformer: a foundation model for molecular generation," Digital Discovery, vol. 4, no. 10, pp. 2684–2696, 2025.
[11] A. Nemecek, W. Li, X. Jiang, J. Vaidya, and E. Ayday, "Quantifying memorization and privacy risks in genomic language models," in Proceedings of the 17th ACM International Conference on Bioinformatics, Computational Biology and Health Informatics. ACM, 2026.
[12] J. Arús-Pous, S. V. Johansson, O. Prykhodko, E. J. Bjerrum, C. Tyrchan, J.-L. Reymond, H. Chen, and O. Engkvist, "Randomized SMILES strings improve the quality of molecular generative models," Journal of Cheminformatics, vol. 11, no. 1, p. 71, 2019.
[13] E. J. Bjerrum, "Smiles enumeration as data augmentation for neural network modeling of molecules," arXiv preprint arXiv:1703.07076, 2017.
[14] M. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang, "Deep learning with differential privacy," in Proceedings of the 2016 ACM SIGSAC conference on computer and communications security, 2016, pp. 308–318.